

Practicle 1:

```
In [7]: import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
```

```
In [8]: # Load the dataset
data = pd.read_csv('Salary_Data.csv')
```

```
In [9]: # Split the data into input (X) and output (y) variables
X = data.iloc[:, :-1].values
y = data.iloc[:, 1].values
```


```
In [10]: # Create a linear regression model
regressor = LinearRegression()
```

```
In [17]: # Train the model
regressor.fit(X, y)
```

```
Out[17]: ▾ LinearRegression
LinearRegression()
```

```
In [12]: # Predict the salaries based on the model
y_pred = regressor.predict(X)
```

```
In [13]: # Visualize the data and regression line
plt.scatter(X, y, color='blue')
plt.plot(X, y_pred, color='red')
plt.title('Simple Linear Regression')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.show()
```



In [14]: *# Print the coefficients and intercept*

```
coefficients = regressor.coef_  
intercept = regressor.intercept_  
print('Coefficients:', coefficients)  
print('Intercept:', intercept)
```

Coefficients: [9449.96232146]

Intercept: 25792.200198668717

Practical 2:

```
In [1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error
```

```
In [2]: # Load the dataset
data = pd.read_csv('Salary_Data.csv')
```

```
In [3]: # Split the dataset into features (X) and target variable (y)
X = data.iloc[:, :-1].values
y = data.iloc[:, -1].values
```

```
In [4]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [5]: # Create a KNN regression model
knn = KNeighborsRegressor(n_neighbors=5)
```

```
In [6]: # Train the model
knn.fit(X_train, y_train)
```

```
Out[6]: KNeighborsRegressor
KNeighborsRegressor()
```

```
In [7]: # Predict the target variable for the test set
y_pred = knn.predict(X_test)
```

```
In [8]: # Calculate the mean squared error
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
```

Mean Squared Error: 50357290.43333333

```
In [9]: # Print the predicted and actual values for the test set
print("Predicted\tActual")
for i in range(len(y_pred)):
    print(f"{y_pred[i]}\t\t{y_test[i]}")
```

Predicted	Actual
115249.0	112635.0
59394.4	67938.0
106311.4	113812.0
71904.8	83088.0
58049.8	64445.0
57499.0	57189.0

Practical 3:

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

```
In [2]: # Load the dataset
data = pd.read_csv('Salary_Data.csv')
```

```
In [3]: # Split the data into features (X) and labels (y)
X = data.iloc[:, :-1].values
y = data.iloc[:, -1].values
```

```
In [4]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [5]: # Create and train the logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)
```

C:\ProgramData\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
Out[5]: * LogisticRegression
LogisticRegression()
```

```
In [6]: # Predict the Labels for the test set
y_pred = model.predict(X_test)
```

```
In [7]: # Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.0

```
In [8]: # Plot the Logistic regression model
plt.scatter(X_train, y_train, color='blue', label='Training Data')
plt.scatter(X_test, y_test, color='green', label='Testing Data')
plt.plot(X_test, model.predict(X_test), color='red', linewidth=2, label='Predicted Salary')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Simple Logistic Regression')
plt.legend()
plt.show()
```



Practical 4:

```
In [10]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

```
In [11]: # Load the dataset
data = pd.read_csv('Salary_Data.csv')
```

```
In [12]: # Split the dataset into features (X) and labels (y)
X = data.iloc[:, :-1].values
y = data.iloc[:, -1].values
```

```
In [4]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [5]: # Create an SVM classifier
classifier = SVC(kernel='linear')
```

```
In [6]: # Train the classifier
classifier.fit(X_train, y_train)
```


```
Out[6]: SVC
SVC(kernel='linear')
```

```
In [7]: # Make predictions on the test set
y_pred = classifier.predict(X_test)
```

```
In [8]: # Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.0

```
In [16]: # Plot the decision boundary
plt.scatter(X[:, 0], y, c=y, cmap='viridis')
ax = plt.gca()
xlim = ax.get_xlim()
ylim = ax.get_ylim()
```



Practical 5:

```
In [11]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

```
In [12]: # Load the dataset
data = pd.read_csv('Salary_Data.csv')
```

```
In [13]: # Split the data into features and target variable
X = data.drop('Salary', axis=1)
y = data['Salary']
```

```
In [14]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [15]: # Initialize the Random Forest classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
In [16]: # Train the classifier
rf_classifier.fit(X_train, y_train)
```

```
Out[16]: ▼      RandomForestClassifier
RandomForestClassifier(random_state=42)
```

```
In [17]: # Predict the target variable for the test set
y_pred = rf_classifier.predict(X_test)
```

```
In [18]: # Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.0

```
In [19]: # Get feature importances
importances = rf_classifier.feature_importances_
features = X.columns
```



```
plt.bar(features, importances)
plt.xlabel('Features')
plt.ylabel('Importance')
plt.title('Feature Importances')
plt.xticks(rotation=45)
plt.show()
```

